

EECS 489

Computer Networks

Winter 2023

Z. Morley Mao

Material with thanks to Aditya Akella, Sugih Jamin, Philip Levis, Sylvia Ratnasamy, Peter Steenkiste, and many other colleagues.

Agenda

- Video streaming
- Datacenter applications

How is video different?

- Often **too large** to send in one GET
- Doesn't even make sense even if it's possible
 - Users may skip forward! \Rightarrow save bandwidth wastage
 - Users' connection quality may change (e.g., switching from WiFi to LTE) \Rightarrow lower resolution to save bandwidth
- **Our focus is on stored video (i.e., not live)**

Why video is important?

- Dominates the global Internet traffic landscape
 - About 60%, i.e., every 3 of 5 bytes in 2020!
- Major sources
 - Netflix
 - YouTube
 - ...

The video medium

- Video is a sequence of images/frames displayed at a constant rate (moving pictures)
- Digital image is an array of pixels, each pixel represented by bits
- Examples:
 - Single frame image encoding: 1024x1024 pixels, 24 bits/pixel \Rightarrow 3 MB/image
 - Movies: 24 frames/sec \Rightarrow 72 MB/sec
 - TV: 30 frames/sec \Rightarrow 90 MB/sec

The video medium (cont'd)

- Compression is key
 - Lots of algorithms to compress
- The same video can be (and typically is) compressed to multiple quality levels
 - E.g., 480p, 720p, 1080p, 4K
- Why multiple resolutions?
 - Adapt to conditions

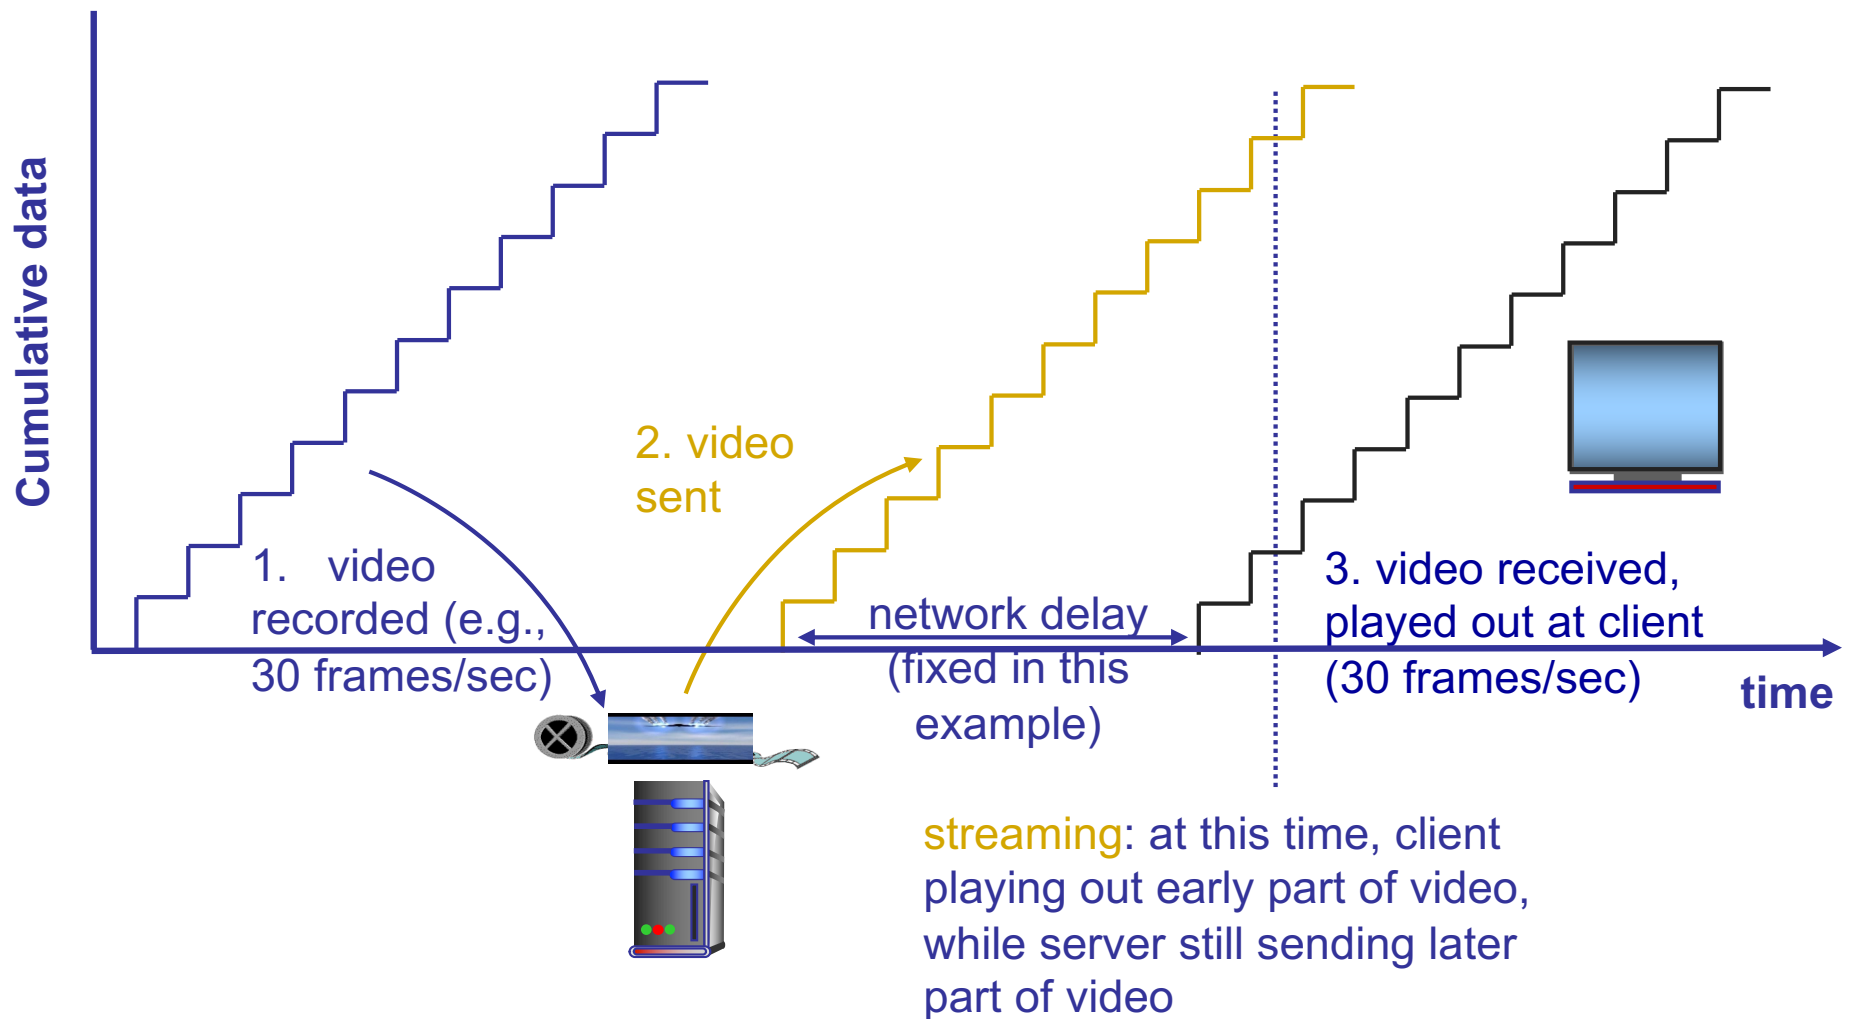
How do we serve video?

- It's in the name!
 - Video streaming

HTTP streaming

- Video is stored at an HTTP server with a URL
- Clients send a GET request for the URL
- Server sends the video file as a stream
- Client first buffers for a while. Why?
 - To minimize interruptions later
- Once the buffer reaches a threshold
 - The video plays in the foreground
 - More frames are downloaded in the background


HTTP streaming



Challenges

- Absorb network delay variations
- Handle user interactions
 - Jump forward, fast-forward, rewind, pause
- Handle packet loss, retransmission etc.

HTTP streaming: Revisited



- **Client-side buffering and playout delay:** compensate for network-added delay, delay jitter

Issues with HTTP streaming


- Same bitrate for all clients
 - Clients can have very different network conditions
 - Clients network conditions can change over time
- Cannot dynamically adapt to conditions

DASH : Dynamic Adaptive Streaming over HTTP


- Keep multiple resolutions of the same video
 - Stored in a manifest file in the HTTP server
- Client asks for the manifest file first to learn about the options
- Asks for chunks at a time and measures available bandwidth while they are downloaded
 - Low bandwidth \Rightarrow switch to lower bitrate
 - High bandwidth \Rightarrow switch to higher bitrate

CLOUD SYSTEMS

Who's serving Web services?



Who's serving Web services?




Cloud datacenters run the world



Source: Google

Cloud datacenters run the world



Source: Google

Cloud datacenters run the world



Source: Facebook

How big is a datacenter (DC)?

- 1M servers/site [Microsoft/Amazon/Google]
- > \$1B to build one site [Facebook]
- >\$20M/month/site operational costs [MS'09]
- Data center hardware spending grew to **\$177 billion** in 2017. [Gartner report]
- But only $O(10-100)$ sites

Implications (1)

- Scale
 - Need scalable designs
 - Low-cost designs: e.g., use commodity technology
 - High utilization (efficiency): e.g., >60% avg. utilization
 - » **Contrast**: avg. utilization on Internet links often ~30%
 - Tolerate frequent failure
 - » Large number of (low cost) components
 - Automate

Implications (2)


- Service model: clouds / multi-tenancy
 - Performance guarantees
 - Isolation guarantees
 - Portability

Applications

- Common theme: **parallelism**
 - Applications decomposed into tasks
 - Running in parallel on different machines
- Two common paradigms
 - Partition-Aggregate
 - Map-Reduce

Partition-Aggregate

eeecs 489



Partition-Aggregate

eeecs 489



GitHub - mosharaf/eeecs489: EECS 489: Computer Networks @ the ...

<https://github.com/mosharaf/eeecs489> ▼

EECS 489: Computer Networks (F'18) ... **EECS 489** takes a top-down approach to explore how networks operate and how network applications are written. ... Kurose and Ross, Computer Networking: A Top-Down Approach, 7th.

UM EECS 489: Computer Networks

www.eecs.umich.edu/courses/eeecs489/ ▼

Lecture: MWF 9:30 - 10:30 in 1500 **EECS**. Discussion/Lab: W 12:30 - 1:30 in 2166 DOW or W 4:30 - 5:30 in 1014 DOW. The discussion sessions will mostly be ...

EECS 489

www.eecs.umich.edu/courses/eeecs489/f99/ ▼

News group umich.**eeecs**.class.**489**. Everything posted here will be automatically forwarded to the **eeecs489staff@eeecs.umich.edu** mailing list hourly.

EECS 489 - EECS @ Michigan - University of Michigan

<https://www.eecs.umich.edu/eeecs/academics/courses/eeecs-489.html> ▼

Course Homepage: <http://www.eecs.umich.edu/courses/eeecs489/w10/>. Coverage We study how networks operate and how network applications are written.

Revamping EECS 489: A Retrospective | Mosharaf Chowdhury

<https://www.mosharaf.com/blog/2017/05/07/revamping-eeecs-489-a-retrospective/> ▼

May 7, 2017 - A couple of weeks ago, we wrapped up the Spring 2017 offering of the **EECS 489**:

End-to-end response time

- Less than 200 milliseconds between receiving user query in the browser and displaying the results
 - $RTT = O(10)$ to 100 milliseconds
 - What remains?
- Next time, when the page is not loading fast enough, think about the poor servers working for you 😊

Announcements


- Midterm is Wed Feb 22 2023
 - 9AM (everyone is expected to take it at this time)
- Assignment 1 due this Fri Jan 27, 2023, 11:59 PM.
- Deadline for forming project groups was yesterday.
- Quiz 5 will be released today (due within 48 hours).

5-MINUTE BREAK!

Applications


- Common theme: parallelism
 - Applications decomposed into tasks
 - Running in parallel on different machines
- Two common paradigms
 - Partition-Aggregate
 - Map-Reduce

Map-Reduce




The most popular software that follows this paradigm is **Apache Spark**


Datacenter networks



Datacenter networks (Cont.)



Datacenter traffic



East-West traffic

- Traffic **between servers** in the datacenter
- Communication within “big data” computations
- Traffic may shift on small timescales (< minutes)


Datacenter traffic characteristics

- Two key characteristics
 - Most flows are small
 - Most bytes come from large flows
- Applications want
 - High bandwidth (large flows)
 - Low latency (small flows)


High bandwidth

- Ideal: Each server can talk to any other server at its full access link rate
- Conceptually: Datacenter network as one giant switch


Datacenter network as one giant switch



Datacenter network as one giant switch



Datacenter network as one giant switch



High bandwidth


- Ideal: Each server can talk to any other server at its full access link rate
- Conceptually: Datacenter network as one giant switch
 - Would require a 10 Pbits/sec switch! (10^{15} bps)
 - » 1M ports (one port/server)
 - » 10Gbps per port
- **Practical approach:** build a network of switches (“fabric”) with high “bisection bandwidth”
 - Each switch has practical #ports and link speeds

Bisection bandwidth


- Partition a network into two equal parts
- Minimum bandwidth between the partitions is the bisection bandwidth
- **Full bisection bandwidth**: bisection bandwidth in an N node network is $N/2$ times the bandwidth of a single link
 - Nodes of any two halves can communicate at full speed with each other

Achieving full bisection bandwidth

- Scale up
 - Make links fatter toward the core of the network
- Problem: Scaling up a traditional tree topology is expensive!
 - Requires non-commodity / impractical / link and switch components
- Solutions?
 - Over-subscribe (i.e., provision less than full BBW)
 - Better topologies



Oversubscription




Need techniques to **avoid congesting oversubscribed links!**


Oversubscription

- Not enough bandwidth
 - **Oversubscription**: Less bandwidth in the ToR-Agg links than all the servers bandwidth in the rack
 - **Oversubscription ratio**: Ratio between bandwidth underneath and bandwidth above
- Not enough paths between server pairs
 - Load balancing issues
 - Failure recovery issues

Better topologies




Better topologies




Clos topology

- Multi-stage network
- k pods, where each pod has two layers of $k/2$ switches
 - $k/2$ ports up and $k/2$ down
- All links have the same b/w
- At most $k^3/4$ machines
- Example
 - $k = 4$
 - 16 machines
- For $k=48$, 27648 machines



Challenges in scale-out designs?

- Topology offers high bisection bandwidth
- All other system components must be able to exploit this available capacity
 - Routing must use all paths
 - Transport protocol must fill all pipes (fast)



Summary

- Video streaming
 - Too large to send, so stream it
 - Dynamically adapt to the network and users
- Cloud systems
 - Forms the backend of modern web services
 - Runs in datacenters where all the processing happens