

# **EECS 489**

# **Computer Networks**

**Winter 2023**

Z. Morley Mao

*Material with thanks to Aditya Akella, Sugih Jamin, Philip Levis, Sylvia Ratnasamy, Peter Steenkiste, and many other colleagues.*

# Logistics

---

- Open book/text/notes, but OFFLINE
  - Except for taking the exam over the Internet
- You're NOT allowed to write/run any programs
- You're NOT allowed to collaborate with anyone

# General guidelines (1)

---

- Test only assumes material covered in lecture, sections, and assignments **after midterm**
  - Text: only to clarify details and context for the above
- The test doesn't require you to do complicated calculations
  - Use this as a hint to determine if you're on right track
- You don't need to memorize anything
- **You do need to understand how things work**

# General guidelines (2)

---

- Be prepared to:
  - Weigh design options outside of the context we studied them in
  - Contemplate new designs we haven't covered in detail but can be put together
    - »e.g., I introduce a new IP address format; how does this affect..”
  - Reason from what you know about the pros/cons of solutions we did study

# General guidelines (3)

---

- Exam format
  - Like midterm, but we're working to avoid cascading mistakes
- Questions not ordered in terms of complexity
  - » Read all carefully
- Pace yourself accordingly!

# This review


---

- Walk through what you're expected to know since the midterm: key topics, important aspects of each
- Not covered in review **does NOT imply** you don't need to know it
  - But if it's covered today, you should know it
- Summarize, not explain
  - Stop me when you want to discuss something further!

# The networking stack

---

- Lower three layers implemented everywhere
- Top two layers implemented only at hosts



# Topics


---

- Network layer (lectures 12–16)
  - Intra-domain routing
  - Inter-domain routing
  - SDN
- Link layer (lectures 17–19)
  - Ethernet
  - Wireless
- Datacenter networking (lectures 20)



# Network layer

- Present everywhere
- Performs **addressing**, **forwarding**, and **routing**, among other tasks

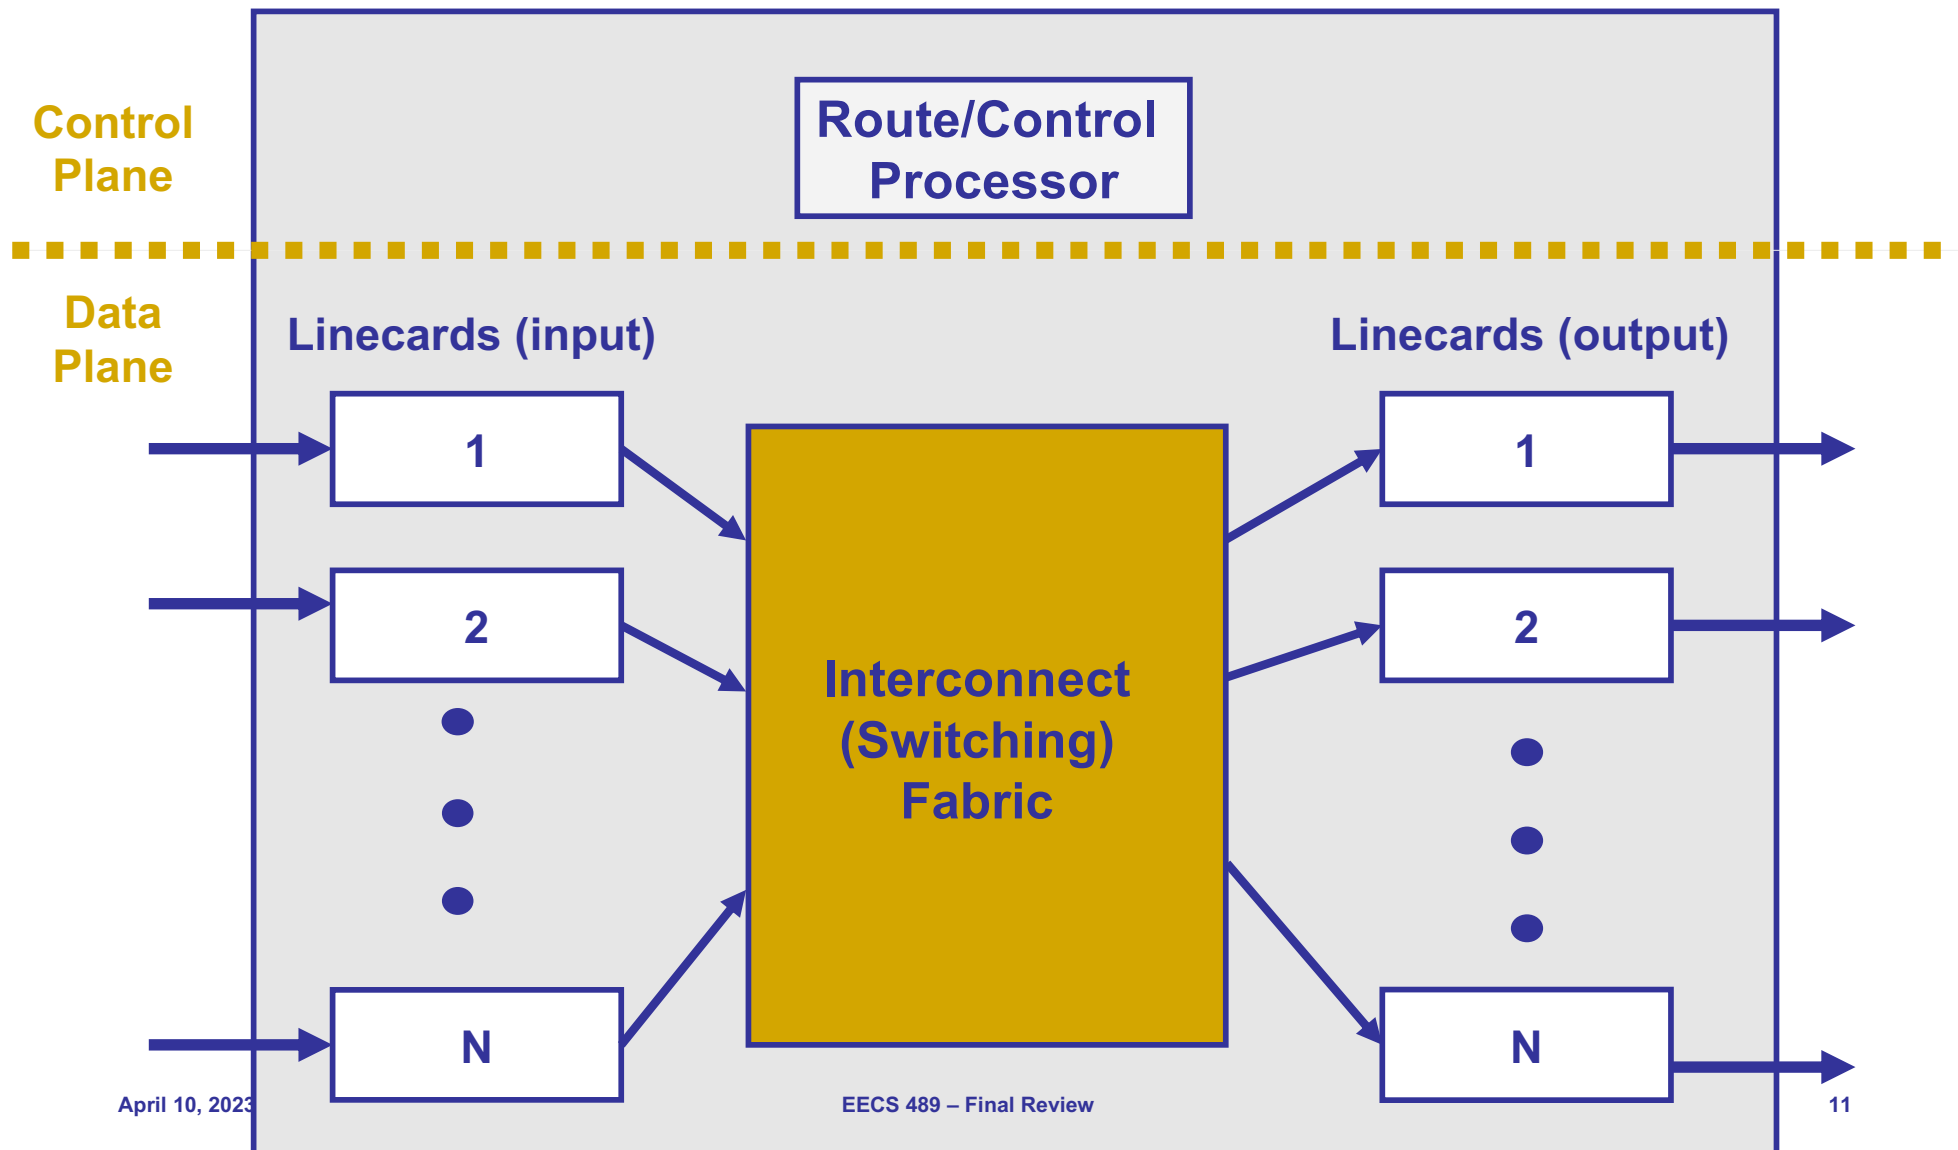


# Forwarding vs. routing

---

- Forwarding: “data plane”
  - Directing one data packet
  - Each router using local routing state
- Routing: “control plane”
  - Computing the forwarding tables that guide packets
  - Jointly computed by routers using a distributed algorithm
- Very different timescales!

# What's inside a router?



# Routing: Local vs. global view

---

- *Local* routing state is the forwarding table in a single router
  - By itself, the state in a single router cannot be evaluated
  - It must be evaluated in terms of the global context
- *Global* state refers to the collection of forwarding tables in each of the routers
  - Global state determines which paths packets take

# “Valid” routing state

---

- Global state is “valid” if it produces forwarding decisions that always deliver packets to their destinations
- Goal of routing protocols: compute valid state
  - How can we tell if routing state is valid?

# Necessary and sufficient condition

---

- Global routing state is valid *if and only if*:
  - There are no dead ends (other than destination)
  - There are no loops
- A **dead end** is when there is no outgoing link (next-hop)
  - A packet arrives, but the forwarding decision does not yield any outgoing link
- A **loop** is when a packet cycles around the same set of nodes forever

# Least-cost routes

---

- Least-cost routes provide an easy way to avoid loops
  - No reasonable cost metric is minimized by traversing a loop
- Least-cost paths form a spanning tree for each destination rooted at that destination

# Intra-domain routing

---

- Link-state (LS) routing protocol
  - Dijkstra's algorithm
  - Broadcast neighbors' info to everyone
- Distance vector (DV) routing protocol
  - Bellman-Ford algorithm
  - Gossip to neighbors about everyone



# Similarities between LS and DV routing

---

- Both are shortest-path based routing
  - Minimizing cost metric (link weights) a common optimization goal
    - » Routers share a common view as to what makes a path “good” and how to measure the “goodness” of a path
- Due to shared goal, commonly used inside an organization
  - RIP and OSPF are mostly used for **intra**-domain routing

# Comparison of LS and DV routing

---

## Messaging complexity

- LS: with  $N$  nodes,  $E$  links,  $O(NE)$  messages sent
- DV: exchange between neighbors only

## Speed of convergence

- LS: relatively fast
- DV: convergence time varies
  - **Count-to-infinity** problem

## Robustness: what happens if router malfunctions?

- LS:
  - Node can advertise incorrect **link** cost
  - Each node computes its *own* table
- DV:
  - Node can advertise incorrect **path** cost
  - Each node's table used by others (**errors propagate**)

# Addressing is key to scalable inter-domain routing

---

- Ability to aggregate addresses is crucial for
  - State: Small forwarding tables at routers
    - » Much less than the number of hosts
  - Churn: Limited rate of change in routing tables

# Classful addressing

---

- Three classes
  - 8-bit network prefix (Class A),
  - 16-bit network prefix (Class B), or
  - 24-bit network prefix (Class C)
- Example: an organization needs 500 addresses.
  - A single class C address is not enough (<500 hosts)
  - Instead, a class B address is allocated (~65K hosts)
    - » Huge waste!


# CIDR: Classless inter-domain routing

---

- Flexible division between network and host addresses
- Offers a better tradeoff between size of the routing table and efficient use of the IP address space

# Hierarchy in IP addressing

- 32 bits are partitioned into a prefix and suffix components
- Prefix is the **network** component; suffix is the **host** component




- Inter-domain routing operates on network prefix

# Administrative structure shapes Inter-domain routing

---

- ASes want freedom to pick routes based on **policy**
- ASes want **autonomy**
- ASes want **privacy**

# Business relationships



## *Relations between ASes*


provider  $\longleftrightarrow$  customer  
peer  $\bullet\text{---}\bullet$  peer

## *Business implications*

- Customers pay provider
- Peers don't pay each other



# Routing follows the money!



- ASes provide “transit” between their customers
- Peers do not provide transit between other peers


# BGP inspired by Distance-Vector with four differences

---


- Shortest-path routes may not be picked to enforce policy
- Path-Vector routing to avoid loops
- Selective route advertisement may affect reachability
- Routes may be aggregated for scalability

# BGP: Basic idea

---



# Policy dictates how routes are “selected” and “exported”



- **Selection:** Which path to use?
  - Controls whether/how traffic leaves the network
- **Export:** Which path to advertise?
  - Controls whether/how traffic enters the network

# Typical export policy

Destination prefix advertised by...	Export route to...
Customer	Everyone (providers, peers, other customers)
Peer	Customers
Provider	Customers

We'll refer to these as the “Gao-Rexford” rules (capture common – **but not required!** – practice)

# Selection using attributes

- Rules for route selection in priority order

Priority	Rule	Remarks
1	LOCAL PREF	Pick highest LOCAL PREF
2	ASPATH	Pick shortest ASPATH length
3	MED	Lowest MED preferred
4	eBGP > iBGP	Did AS learn route via eBGP (preferred) or iBGP?
5	iBGP path	Lowest IGP cost to next hop (egress router)
6	Router ID	Smallest next-hop router's IP address as tie-breaker

# eBGP, iBGP, and IGP

---

- **eBGP**: BGP sessions between border routers in different ASes
  - Learn routes to external destinations
- **iBGP**: BGP sessions between border routers and other routers within the same AS
  - Distribute externally learned routes internally
- **IGP**: “Interior Gateway Protocol” = Intra-domain routing protocol
  - Provide internal reachability via shortest path
  - E.g., OSPF, RIP

---

**5-MINUTE BREAK!**



# Announcements

---

- Final exam date and time:
  - 90 Min; Online (same setup as Midterm exam)
  - Friday April 21: 1:30 pm - 3:00 pm
  - From registrar's final exam schedule
  - For those needing extra time: same start time (longer exam time)

# “The Power of Abstraction”

---

- “Modularity based on abstraction is the way things get done”
  - Barbara Liskov
- Abstractions → Interfaces → Modularity


# Separate concerns with abstractions

---

- Be compatible with low-level hardware/software
  - Need an abstraction for general forwarding model
- Make decisions based on entire network
  - Need an abstraction for network state
- Compute configuration of each physical device
  - Need an abstraction that simplifies configuration

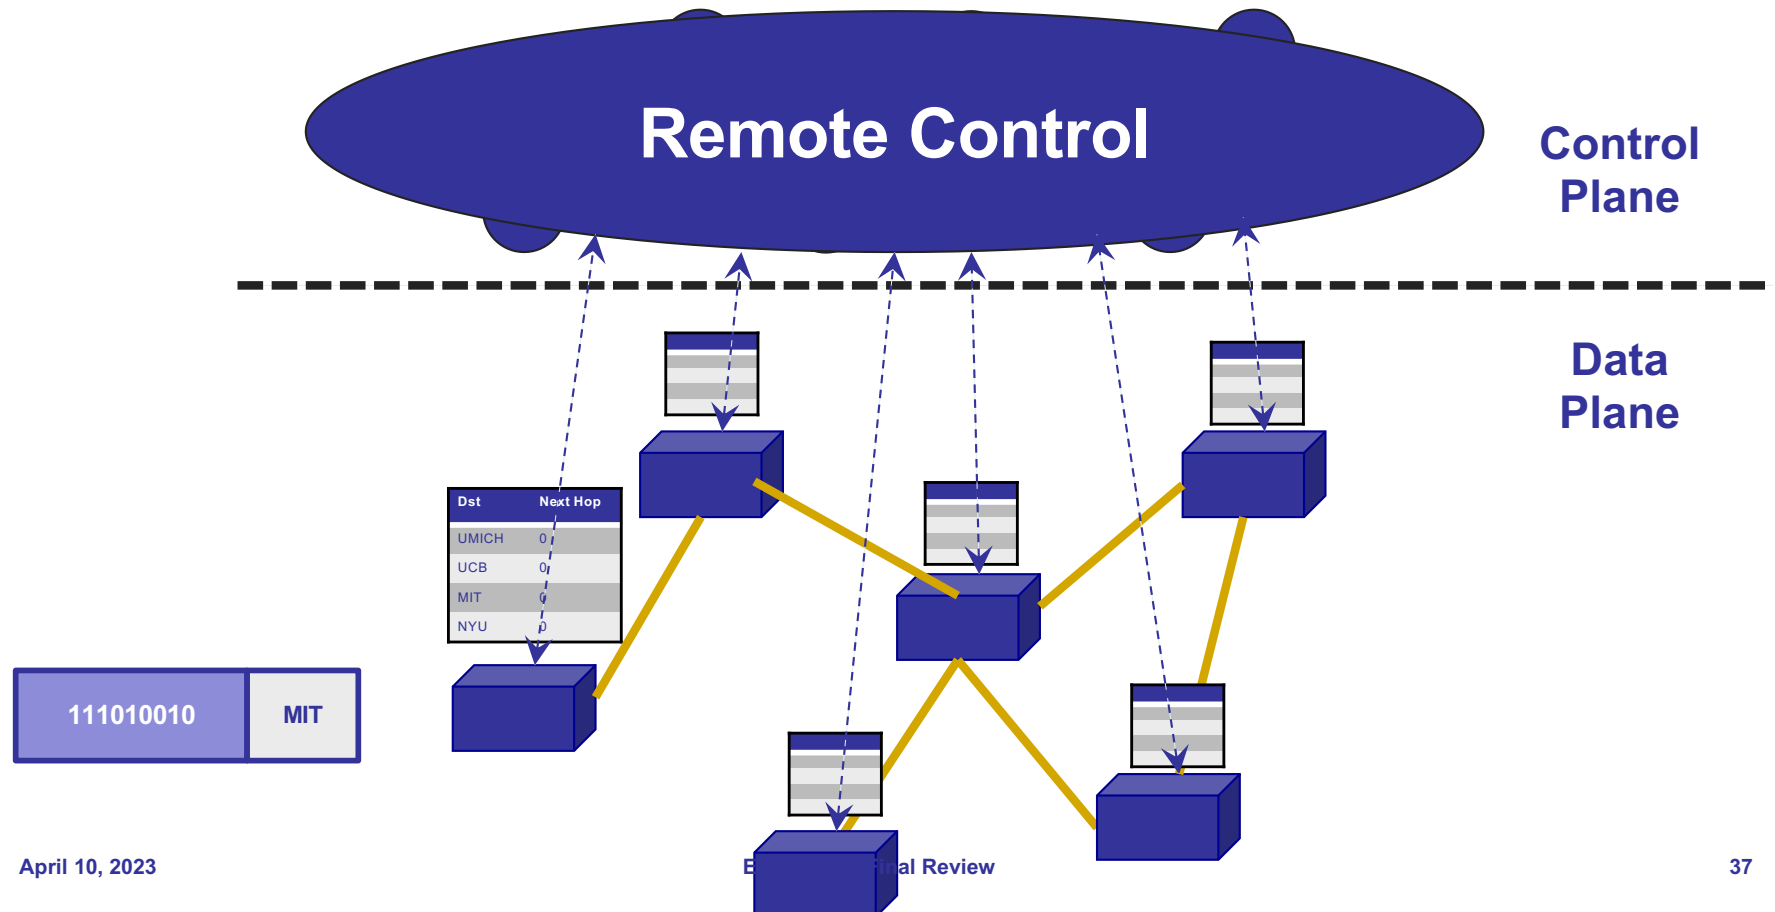
# Traditional fully decentralized control plane

- Individual routing algorithm components in every router interact in the control plane



# Logically centralized control plane

- A distinct (typically remote) controller interacts with local control agents (CAs)



# SDN: Many challenges remain

---

- Hardening the control plane: dependable, reliable, performance-scalable, secure distributed system
  - **Robustness to failures**: leverage strong theory of reliable distributed system for control plane
  - **Security**: “baked in” from day one?
- Networks, protocols meeting mission-specific requirements
  - E.g., real-time, ultra-reliable, ultra-secure
- **Internet-scaling**

# Fixed-function data plane

---

- Traditional switches are fixed-function
  - They can do whatever they can do at birth, but they cannot change!
  - Bottom-up design
- Even OpenFlow was designed to be a fixed protocol
  - With a fixed table format
  - Capable of doing limited things

# Programmable data plane

---

- What if we could tell switches exactly what we want?
  - What table to keep?
  - What rules to use?
  - What data to keep track of?
  - ...

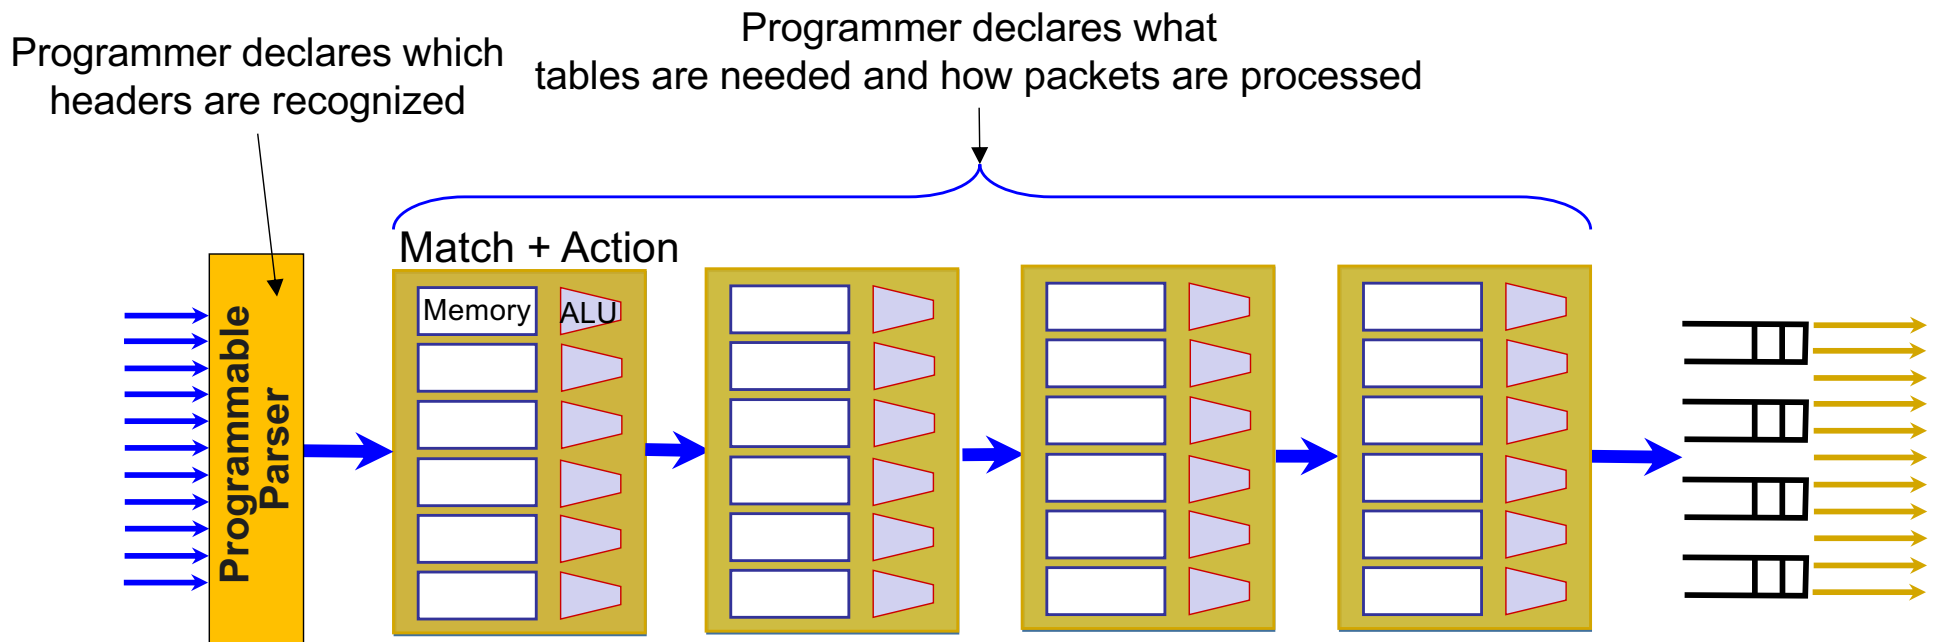


# Top-down workflow

---

- Precisely specify using a well-defined language
- Compile it down to run on a standardized hardware (e.g., using P4)
- Run at line speed

# PISA: Protocol Independent Switch Architecture



**All stages are identical – makes PISA a good “compiler target”**

# Topics

---

- Network layer (lectures 12–16)
  - Intra-domain routing
  - Inter-domain routing
  - SDN
- Link layer (lectures 17–19)
  - Ethernet
  - Wireless
- Datacenter networking (lectures 20)

# Data link layer

---

- Provides four primary services
  - Framing
    - » Encapsulates network layer data
  - Link access
    - » Medium access control (MAC) protocol defines when to transmit frames
  - Reliable delivery
    - » Primarily for mediums with high error rates (e.g., wireless)
  - Error detection and correction

# Point-to-point vs. broadcast medium

---

- **Point-to-point**: dedicated pairwise communication
  - E.g., long-distance fiber link
  - E.g., Point-to-point link b/n Ethernet switch and host
- **Broadcast**: shared wire or medium
  - Traditional Ethernet (pre ~2000)
  - 802.11 wireless LAN

# Random access MAC protocols

---

- When node has packet to send
  - Transmit at full channel data rate **w/o** coordination
- Two or more transmitting nodes  $\Rightarrow$  **collision**
  - Data lost
- Random access MAC protocol specifies
  - How to **detect** and **recover** from collisions
- Examples
  - ALOHA and Slotted ALOHA
  - **CSMA**, **CSMA/CD**, CSMA/CA (wireless)

# CSMA (Carrier Sense Multiple Access)

---

- CSMA: listen before transmit
  - If channel sensed idle: transmit entire frame
  - If channel sensed busy, defer transmission
- Human analogy: don't interrupt others!
- Does not eliminate all collisions
  - Why?

# CSMA/CD (Collision Detection)

---

- CSMA/CD: carrier sensing, deferral as in CSMA
  - Collisions detected within short time
  - Colliding transmissions aborted, reducing wastage
- Collision detection easy in wired (broadcast) LANs
  - Compare transmitted, received signals
- Collision detection difficult in wireless LANs



# Limits on CSMA/CD network length



- Latency depends on physical length of link
  - Time to propagate a frame from one end to other
- Suppose A sends a frame at time  $t$ 
  - And B sees an idle line at a time just before  $t + d$
  - ... so B happily starts transmitting a frame
- B detects a collision, and sends jamming signal
  - But A cannot see collision until  $t + 2d$


# Limits on CSMA/CD network length



- A needs to wait for time  **$2d$**  to detect collision
  - So, A should keep transmitting during this period
  - AND keep an eye out for a possible collision
- Imposes restrictions; e.g., for 10 Mbps Ethernet
  - **Maximum length** of the wire: 2,500 meters
  - **Minimum length** of a frame: 512 bits (64 bytes)

# Why switched Ethernet?


- Enables concurrent communication
  - Host A can talk to C, while B talks to D
  - No collisions and no need for CSMA/CD
  - No constraints on link lengths, etc.



# Ethernet switches are “self learning”

- When a packet arrives:
  - Inspect source MAC address, associate with incoming port
  - Store mapping in the switch table
  - Use **time-to-live** field to eventually forget mapping

Packet tells switch  
how to reach A.



# ARP and DHCP

---

- Link layer discovery protocols
  - ARP → Address Resolution Protocol
  - DHCP → Dynamic Host Configuration Protocol
  - Confined to a single local-area network (LAN)
  - Rely on broadcast capability

# Key ideas in both ARP and DHCP

---

- **Broadcasting**: Can use broadcast to make contact
  - Scalable because of limited size
- **Caching**: remember the past for a while
  - Store the information you learn to reduce overhead
- **Soft state**: eventually forget the past
  - Associate a time-to-live field with the information
  - ... and either refresh or discard the information
  - Key for robustness in the face of unpredictable change


# ARP: Address Resolution Protocol

---

- Every host maintains an ARP table
  - List of (IP address → MAC address) pairs
- Consult the table when sending a packet
  - Map dest. IP address to dest. MAC address
  - Encapsulate (IP) data packet with MAC header; xmit
- What if IP address not in the table?
  - Sender broadcasts: Who has IP address 1.2.3.156?
  - Receiver replies: MAC address 58-23-D7-FA-20-B0
  - Sender caches result in its ARP table

# What if the destination is remote?

- Look up the MAC address of the first hop router
  - 1.2.3.48 uses ARP to find MAC address for first-hop router **1.2.3.19** rather than ultimate destination IP address
- How does the red host know the destination is not local?
  - Uses **netmask** (discovered via DHCP)
- How does the red host know about 1.2.3.19?
  - Also DHCP





# Wireless link characteristics

---

- Three important differences from wired link ...
  - **Decreased signal strength**: Radio signal attenuates as it propagates through matter (path loss)
  - **Multipath propagation**: Radio signal reflects off objects ground, arriving at destination at slightly different times
  - **Interference from other sources**: Standardized wireless network frequencies (e.g., 2.4 GHz) shared by other devices (e.g., phone); devices (motors) interfere as well
- ... make communication across (even a point-to-point) wireless link much more “difficult”


# Wireless network characteristics

---


- Multiple wireless senders and receivers create many problems
  - Multiple access issues (we've seen this before)
  - Hidden terminal problem

# Hidden terminal problem

- B, A hear each other
- B, C hear each other
- A, C can not hear each other
- Hence, A, C are unaware of their interference at B



# CSMA/CA



- Before every data transmission
  - Sender sends a Request to Send (RTS) frame with the length of transmission and the destination
  - Receiver respond with a Clear to Send (CTS) frame
  - Sender sends data
  - Receiver sends an ACK
- If sender doesn't get a CTS back, it assumes collision

# Topics

---

- Network layer (lectures 12–16)
  - Intra-domain routing
  - Inter-domain routing
  - SDN
- Link layer (lectures 17–19)
  - Ethernet
  - Wireless
- Datacenter networking (lectures 20)

# Datacenter applications

---

- Common theme: **parallelism**
  - Applications decomposed into tasks
  - Running in parallel on different machines
- Two common paradigms
  - Partition-Aggregate
  - Map-Reduce


# Datacenter traffic characteristics

---

- Two key characteristics
  - Most flows are small
  - Most bytes come from large flows
- Applications want
  - High bandwidth (large flows)
  - Low latency (small flows)

# Clos topology

- Multi-stage network
- $k$  pods, where each pod has two layers of  $k/2$  switches
  - $k/2$  ports up and  $k/2$  down
- All links have the same b/w
- At most  $k^3/4$  machines
- Example
  - $k = 4$
  - 16 machines
- For  $k=48$ , 27648 machines



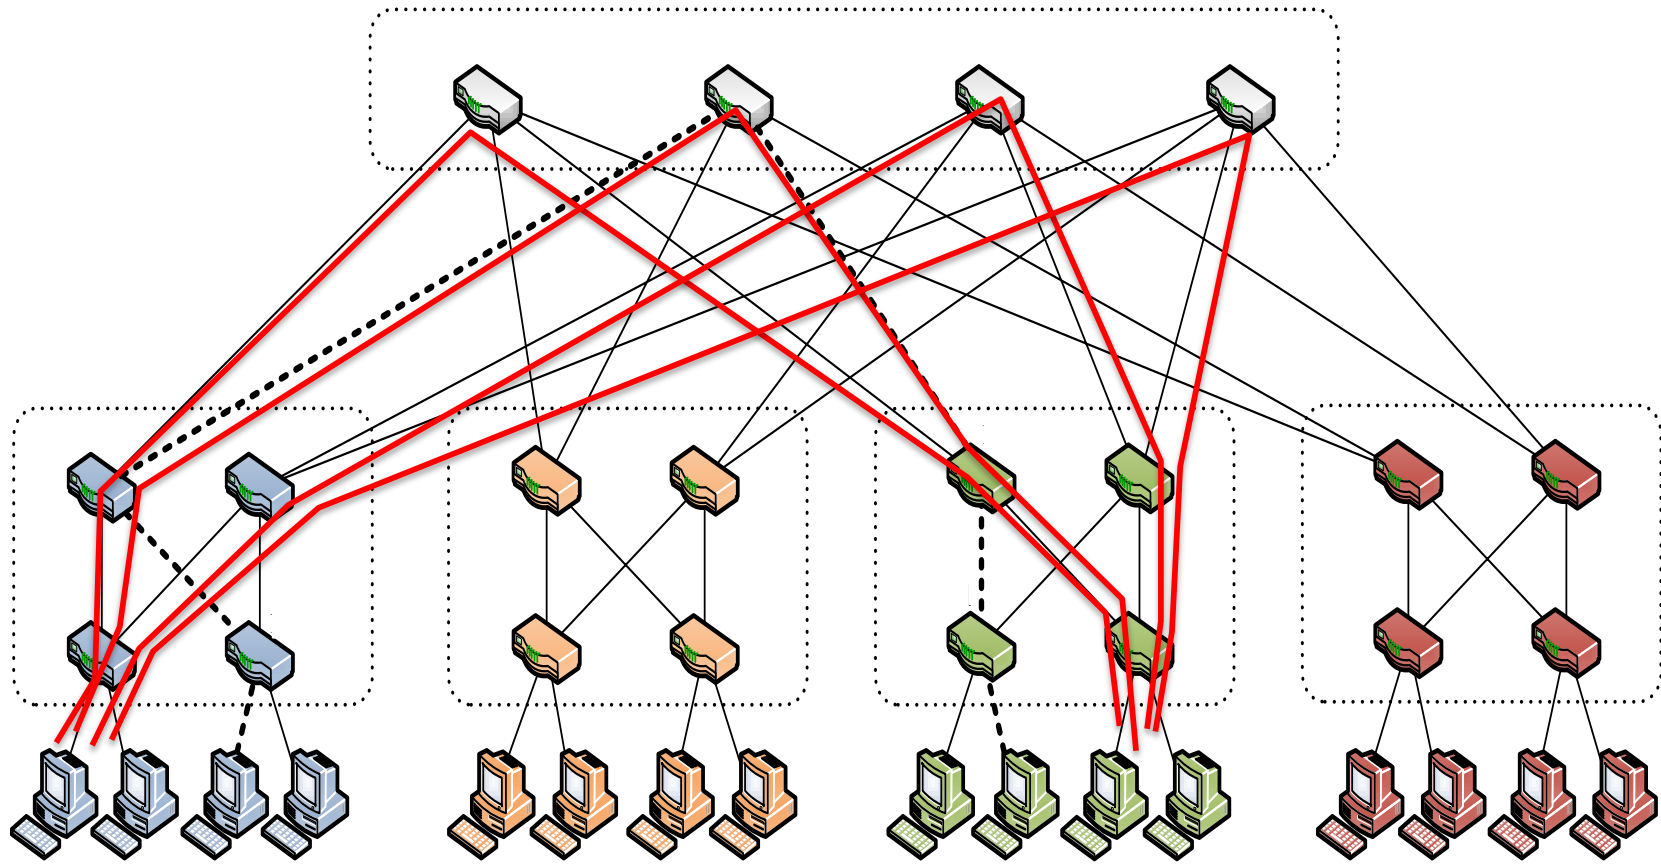


# Datacenter networking stack

---

- Networking in modern datacenters
  - L2/L3 design
    - » Addressing / routing / forwarding in the Fat-Tree
  - L4 design
    - » Transport protocol design (w/ Fat-Tree)
  - L7 design
    - » Exploiting application-level information (w/ Fat-Tree)

# Using multiple paths well



# L2/L3 highlights

---

- Load balancing while forwarding
  - Per-packet
  - Per-flow
- Hard-coded addressing or via indirection
- Modified LS/DV or source routing

# L4 highlights

---

- Tension between high throughput and low latency requirements
  - Deep queues vs shallow queues
- DCTCP
  - React early, quickly, and with certainty using ECN
  - React in **proportion to the extent of congestion**, not its presence

# L7 highlights

---

- What do applications care about?
  - Flow completion time (FCT)
  - Coflow completion time (CCT)
    - » A coflow is a collection of flows with a shared application-level objective
  - We should strive to optimize as close an objective as possible to the application

# Summary

---

- THANK YOU SO MUCH!!!